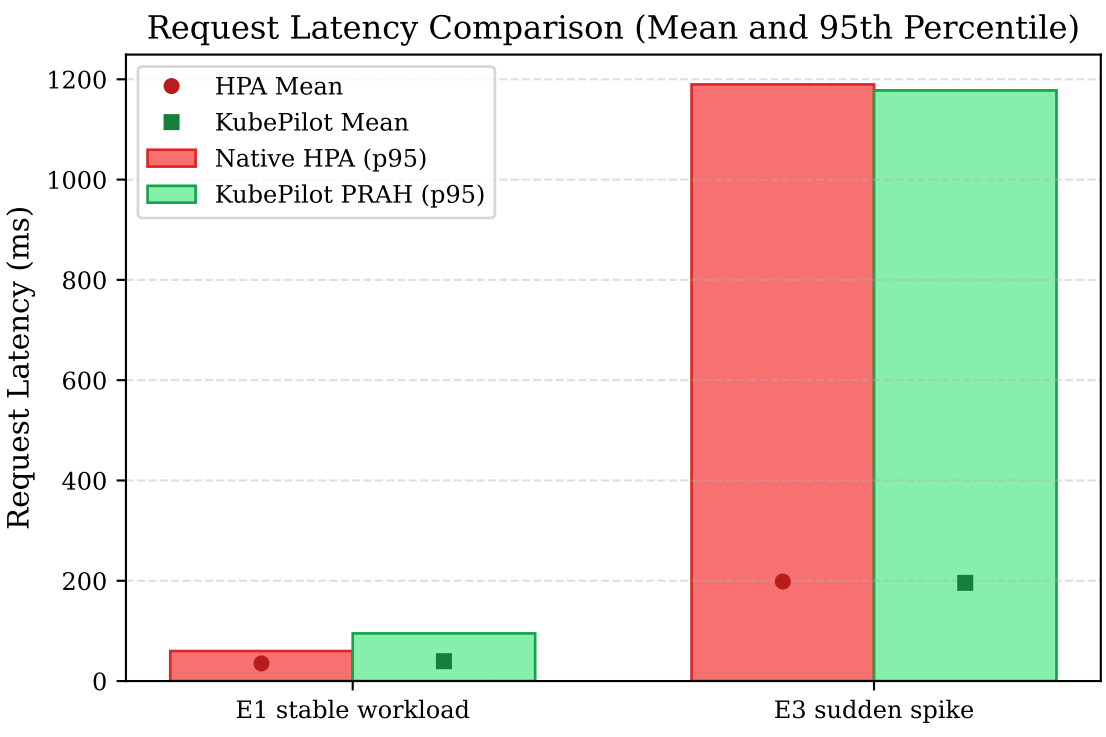
Request Latency Comparison (Mean and 95th Percentile)
Request Latency (ms)
0
200
400
600
800
1000
1200
HPA Mean
KubePilot Mean
Native HPA (p95)
KubePilot PRAH (p95)
E1 stable workload
E3 sudden spike