Request Latency Comparison Across Evaluated Scenarios (Mean and 95th Percentile)
Request Latency (ms)
2000
1750
1500
1250
1000
750
500
250
0
HPA Mean
KubePilot Mean
Native HPA (p95)
KubePilot PRAH (p95)
E1 stable workload
E2 gradual increase
E3 sudden spike
E4 bursty workload
E5 sudden decrease